Package 'MuMIn'

October 15, 2011

Type Package

Title Multi-model inference

Version 1.5.0

Date 2011-10-13

Encoding UTF-8

Author Kamil Bartoń

Maintainer Kamil Bartoń <kamil.barton@go2.pl>

Description Model selection and model averaging based on information criteria (AICc and alike).

License GPL-2

Depends R (>= 2.12.0)

Imports stats

Suggests stats4, nlme, mgcv (>= 1.7.5), lme4 (>= 0.999375-
16), gamm4,MASS, nnet, spdep, survival, unmarked, glmmML

LazyLoad yes

R topics documented:

```
MuMIn-package          Multi-model inference
```

Description

The package `MuMIn` contains functions to streamline model selection and perform model averaging based on information criteria (AIC, AICc and alike).

Details

User level functions include:

dredge performs automated model selection based on subsets of the supplied 'global' model, and optional list of choices for other model characteristics (e.g. different link functions). Model set may be generated either with "all possible" combinations, or tailored according to complex inclusion rules

model.sel creates a model selection table from handpicked models

model.avg calculates model averaged parameters

AICc calculates second-order Akaike information criterion for one or several fitted model objects.

Model selection can be done according to any information criterion, including AIC, AICc, BIC, QAIC, ICOMP or Mallows' Cp.

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

See Also

AIC, step

Examples

```
fm1 <- lm(Fertility ~ . , data = swiss)

dd <- dredge(fm1)
top.models.1 <- get.models(dd, subset = delta < 4)
model.avg(top.models.1) # get averaged coefficients

top.models.2 <- get.models(dd, cumsum(weight) <= .95)
model.avg(top.models.2)

# Mixed models:
# modified example(lme)
data(Orthodont, package="nlme")
require(nlme)
```

```
fm2 <- lme(distance ~ age + Sex, data = Orthodont,
    random = ~ 1 | Subject, method="ML")
dredge(fm2)
```

AICc   *Second-order Akaike Information Criterion*

Description

Calculates second-order Akaike information criterion for one or several fitted model objects (AIC for small samples).

Usage

```
AICc(object, ..., k = 2, REML = NULL)
```

Arguments

object	a fitted model object for which there exists a `logLik` method, or a `logLik` object
...	optionally more fitted model objects
k	the "penalty" per parameter to be used; the default `k = 2` is the classical `AIC`
REML	optional logical value, passed to the `logLik` method indicating whether the restricted log-likelihood or log-likelihood should be used. The default is to use the method used for model estimation.

Value

If just one object is provided, returns a numeric value with the corresponding AICc; if more than one object are provided, returns a `data.frame` with rows corresponding to the objects and columns representing the number of parameters in the model (df) and AICc.

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

See Also

Akaike's An Information Criterion: `AIC`

`AICc` in package **AICcmodavg**, `aicc` in package **glmulti**

Examples

```
#Model-averaging mixed models

library(nlme)
data(Orthodont, package = "nlme")

# Fit model by REML
fm2 <- lme(distance ~ Sex*age, data = Orthodont,
    random = ~ 1|Subject / Sex, method = "REML")

# Model selection: ranking by AICc using ML
dd <- dredge(fm2, trace=TRUE, rank="AICc", REML=FALSE)

(attr(dd, "rank.call"))

# Get the models (fitted by REML, as in the global model)
gm <- get.models(dd, 1:4)

# Because the models originate from 'dredge(..., rank=AICc, REML=FALSE)',
# the default weights in 'model.avg' are ML based:
model.avg(gm, method = "NA")

# same result
#model.avg(gm, method = "NA", rank="AICc", rank.args = list(REML=FALSE))
# REML based weights
model.avg(gm, method = "NA", rank="AICc", rank.args = list(REML=TRUE))
```

Beetle   *Flour beetle mortality data*

Description

Mortality of flour beetles (*Tribolium confusu*) due to exposure to gaseous carbon disulfide CS_2, from Bliss (1935)

Usage

```
data(Beetle)
```

Format

Beetle is a data frame with 4 variables.

dose The dose of CS_2 in mg/L

n.tested Number of beetles tested

n.killed Number of beetles killed

mortality Observed mortality rate

Source

Bliss C. I. (1935) The calculation of the dosage-mortality curve. *Annals of Applied Biology*, 22: 134–167

References

Burnham, K. P. and Anderson, D. R. (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

Examples

```
# The "Logistic regression example"
# from Burnham & Anderson (2002) chapter 4.11

data(Beetle)

# Fit a global model with all the considered variables
globmod <- glm(cbind(n.killed, n.tested- n.killed)~dose + I(dose^2) + log(dose)
+ I(log(dose)^2), data=Beetle, family=binomial)


# A logical expression defining the subset of models to use:
# * either log(dose) or dose
# * the quadratic terms can appear only together with linear terms
msubset <- expression(xor(dose, `log(dose)`) & (dose | !`I(dose^2)`)
& (`log(dose)` | !`I(log(dose)^2)`))


# Table 4.6

# Use 'varying' argument to fit models with different link functions
# Note the use of 'alist' rather than 'list' in order to keep the 'family'
# objects unevaluated
varying.link <- list(family=alist(
logit = binomial("logit"),
probit = binomial("probit"),
cloglog = binomial("cloglog")
))

(dd12 <- dredge(globmod, subset = msubset, varying = varying.link, rank=AIC))

# Table 4.7 "models justifiable a priori"
(dd3 <- dredge(update(globmod, . ~ dose), m.min = 1, rank=AIC,
varying = varying.link))

mod3 <- get.models(dd3, 1:3)

# Table 4.8. Predicted mortality probability at dose 40.

# helper function to calculate confidence intervals on logit scale
logit.ci <- function(p, se, quantile = 2) {
C. <- exp(quantile * se / (p * (1 - p)))
p /(p + (1 - p) * c(C., 1/C.))
}
```

```
pred <- sapply(mod3, predict, newdata=list(dose=40), se.fit=TRUE, type="response")
pred <- apply(pred, 1, unlist)[,1:2] # simplify

# build the table
tab <- rbind(pred, par.avg(pred[,"fit"], pred[,"se.fit"], dd3$weight,
revised.var = FALSE)[1:2])
tab <- cbind(
c(dd3$weight, NA),
tab,
matrix(logit.ci(tab[,"fit"], tab[,"se.fit"], quantile = c(rep(1.96, 3), 2)),
ncol=2)
)
colnames(tab) <- c("Akaike weight", "Predicted(40)", "SE", "Lower CI", "Upper CI")
rownames(tab) <- c(as.character(dd3$family), "model averaged")

print(tab, digits=3, na.print="")


# Figure 4.3
newdata <- list(dose = seq(min(Beetle$dose), max(Beetle$dose), length.out = 25))
matplot(newdata$dose, sapply(mod3, predict, newdata, type="response"),
type="l", xlab=quote(list("Dose of"~ CS[2],(mg/L))),
ylab="Mortality"
)
```

Cement	*Cement hardening data*

Description

Cement hardening data from Woods et al (1939).

Usage

```
data(Cement)
```

Format

Cement is a data frame with 5 variables. x1-x4 are four predictor variables expressed as a percentage of weight.

X1 calcium aluminate

X2 tricalcium silicate

X3 tetracalcium alumino ferrite

X4 dicalcium silicate

y calories of heat evolved per gram of cement after 180 days of hardening

Source

Woods H., Steinour H.H., Starke H.R. (1932) Effect of composition of Portland cement on heat evolved during hardening. *Industrial & Engineering Chemistry* 24, 1207-1214

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

---

dredge	*Evaluate "all possible" models*

Description

Automatically generate models with combinations of the terms in the global model, with optional rules for inclusion.

Usage

```
dredge(global.model, beta = FALSE, evaluate = TRUE, rank = "AICc",
    fixed = NULL, m.max = NA, m.min = 0, subset, marg.ex = NULL,
    trace = FALSE, varying = NULL, ...)

## S3 method for class 'model.selection'
print(x, abbrev.names = TRUE, ...)
```

Arguments

global.model	a fitted 'global' model object. See 'Details' for a list of supported types.
beta	logical, should standardized coefficients be returned?
evaluate	whether to evaluate and rank the models. If FALSE, a list of model calls is returned.
rank	optional custom rank function (information criterion) to be used instead AICc, e.g. QAIC or BIC. See 'Details'.
fixed	optional, either a single sided formula or a character vector giving names of terms to be included in all models.
m.max, m.min	optional, maximum and minimum number of terms in a single model (excluding the intercept), m.max defaults to the number of terms in global.model.
subset	logical expression to put constraints for the set of models. Can contain any of the global.model terms (use getAllTerms(global.model) to list them). It can have a form of an unevaluated call, expression object, or a one sided formula.
	Compound model terms (such as 'as-is' expressions within I or smooths in gam) should be treated as non-syntactic names and enclosed in back-ticks (see Quotes). Mind the spacing, names must match exactly the term names in model's formula. To simply keep certain variables in all models, use of fixed is preferred.
marg.ex	a character vector specifying names of variables for which NOT to check for marginality restrictions when generating model formulas. If this argument is set to TRUE, all model formulas are used (i.e. no checking). See 'Details'.
trace	if TRUE, all calls to the fitting function (i.e. updated global.model calls) are printed.

varying	optionally, a named list describing the additional arguments to vary between the generated models. Names are the names of the arguments, and each item provides a list of choices. Complex items in the choice list (such as `family` objects) should be either named (uniquely) or quoted (unevaluated, e.g. using `alist`, see `quote`), otherwise it may produce rather unpleasant effects. See example in `Beetle`.
x	a `model.selection` object, returned by `dredge`.
abbrev.names	should variable names be abbreviated when printing? (useful with many variables).
...	optional arguments for the `rank` function. Any can be an expression (of mode `call`), in which case any `x` within it will be substituted with a current model.

Details

`dregde` currently is known to work with `lm`, `glm`, `rlm`, `polr`, `multinom`, `gam`, `gls`, `lme`, `lmer`, `coxph`, `glmmML`, `sarlm`, `spautolm`, `gamm` and `gamm4` (the last two should be fitted *via* the wrapper `MuMIn::gamm`).

Models are run one by one by repeated evaluation of the call to `global.model` with modified formula argument (or `fixed` in `lme`). This method, while robust in that it can be applied to a variety of different models is not very efficient and may be considerably time-intensive.

Note that the number of combinations grows exponentially with number of predictor variables (2^N). Because there is potentially a large number of models to evaluate, to avoid memory overflow the fitted model objects are not stored. To get (a subset of) the models, use `get.models` with the object returned by `dredge` as an argument.

Handling interactions, `dredge` respects marginality constraints, so "all possible combinations" do not include models containing interactions without their respective main effects. This behaviour can be altered by `marg.ex` argument. It can be used to allow for simple nested designs. For example, with global model of form `a / (x + z)`, use `marg.ex = "a"` and `fixed = "a"`.

`rank` is found by a call to `match.fun` and may be specified as a function or a symbol (e.g. a back-quoted name) or a character string specifying a function to be searched for from the environment of the call to `dredge`.

Function `rank` must be able to accept model as a first argument and must always return a scalar. Typical choice for `rank` would be "AIC", "QAIC" or "BIC" (**stats** or **nlme**).

Use of `na.action = na.omit` (R's default) in `global.model` should be avoided, as it results with sub-models fitted to different data sets, if there are missing values. In versions >= 0.13.17 a warning is given in such a case.

Value

`dredge` returns an object of class `model.selection`, being a `data.frame` with models' coefficients (or `TRUE`/`FALSE` for factors), k, deviance/RSS, R-squared, AIC, AICc, delta and weight. This depends on a type of model. Models are ordered according to the used information criterion (lowest on top), specified by `rank`.

The attribute `"calls"` is a list containing the model calls used (arranged in the same order as the models).

Note

Users should keep in mind the hazards that a "thoughtless approach" of evaluating all possible models poses. Although this procedure is in certain cases useful and justified, it may result in selecting a spurious "best" model, due to model selection bias.

"Let the computer find out" is a poor strategy and usually reflects the fact that the researcher did not bother to think clearly about the problem of interest and its scientific setting (Burnham and Anderson, 2002).

Author(s)

Kamil Bartoń

See Also

get.models, model.avg. QAIC has examples of using custom rank function.

There is also subset.model.selection method.

Consider the alternatives: glmulti in package **glmulti** and bestglm (**bestglm**), or aictab (**AICcmodavg**) and ICtab (**bbmle**) for "hand-picked" model selection tables.

Examples

```
# Example from Burnham and Anderson (2002), page 100:
data(Cement)
lm1 <- lm(y ~ ., data = Cement)
dd <- dredge(lm1)
subset(dd, delta < 4)

#models with delta.aicc < 4
model.avg(get.models(dd, subset = delta < 4)) # get averaged coefficients

#or as a 95% confidence set:
top.models <- get.models(dd, cumsum(weight) <= .95)

model.avg(top.models) # get averaged coefficients

#topmost model:
top.models[[1]]

## Not run:
# Examples of using 'subset':
# exclude models containing both X1 and X2
dredge(lm1, subset = !(X1 & X2))
# keep only models containing X3
dredge(lm1, subset = X3)
# the same, but more effective:
dredge(lm1, fixed = "X3")

#Reduce the number of generated models, by including only those with
# up to 2 terms (and intercept)
dredge(lm1, m.max = 2)


## End(Not run)
```

---

gamm   *GAMM wrapper for use with MuMIn*

Description

Adapter function for use mixed Generalized Additive Models from packages **mgcv** and **gamm** with **MuMIn**

Usage

```
gamm(formula, random = NULL, ..., lme4 = inherits(random, "formula"))
```

Arguments

```
formula, random, ...
```
                arguments passed to `gamm` or `gamm4`

```
lme4
```
                logical, if true `gamm4` is used rather than `gamm`, If `TRUE`, the `random` argument must be provided as a `formula`.

Value

Depending on the value of the 'lme4' switch, either a `gamm` or `gamm4` fitted model object. The only difference from the originals is the `call` component, allowing for [update](#)'ing of the object.

Note

This is only a temporary workaround and it is likely be removed soon.

To make sure this wrapper function is used (in case it is masked by the original `gamm` from **mgcv**), use `MuMIn::`[gamm](#).

Author(s)

Kamil Bartoń

See Also

[gamm](#) and [gamm4](#)

---

get.models   *Get models*

Description

Gets list of models from a `model.selection` object

Usage

```
get.models(dd, subset = delta <= 4, ...)
```

Arguments

`dd`	object returned by `dredge`
`subset`	subset of models
`...`	additional parameters passed to `update`, for example, in `lme`/`lmer` one may want to use `method = "REML"` while using "ML" for model selection

Value

`list` of models.

Author(s)

Kamil Bartoń

See Also

`dredge`, `model.avg`

Examples

```
# Mixed models:

require(nlme)
fm2 <- lme(distance ~ age + Sex, data = Orthodont,
    random = ~ 1 | Subject, method="ML")
dd2 <- dredge(fm2)

# Get top-most models, but fitted by REML:
(top.models.2 <- get.models(dd2, subset = delta < 4, method = "REML"))
```

---

importance	*Relative variable importance*

Description

Sum of 'Akaike weights' over all models that include the explanatory variable.

Usage

```
importance(x)
```

Arguments

x	Either a model list or a `"model.selection"` or `"averaging"` object

Value

a numeric vector of relative importance values, named as the variables

Author(s)

Kamil Bartoń

See Also

dredge, model.avg, mod.sel

Examples

```
# Generate some models
data(Cement)
lm1 <- lm(y ~ ., data = Cement)
dd <- dredge(lm1)

# Importance can be calculated/extracted from various objects:
importance(dd)
## Not run:
importance(subset(mod.sel(dd), delta <= 4))
importance(model.avg(dd, subset = delta <= 4))
importance(subset(dd, delta <= 4))
importance(get.models(dd, delta <= 4))

## End(Not run)



# Re-evaluate the importances according to BIC
# note that re-ranking involves fitting the models again

# 'nobs' is not used here for backwards compatibility
lognobs <- log(length(resid(lm1)))

importance(subset(mod.sel(dd, rank=AIC, rank.args=list(k = lognobs)),
cumsum(weight) <= .95))

# This gives a different result than previous command, because 'subset' is
# applied to the original selection table that is ranked with 'AICc'
importance(model.avg(dd, rank=AIC, rank.args=list(k = lognobs),
subset=cumsum(weight) <= .95))
```

Information criteria

Various information criteria

Description

These functions implement Mallows' *Cp* and ICOMP information criterion by Bozdogan (1988)

Usage

```
Cp(object, dispersion = NULL)
ICOMP(object, ..., REML = NULL)
```

Arguments

`object`	a fitted model object (in case of ICOMP, a `logLik` method must exist for the object)
`...`	optionally more fitted model objects
`dispersion`	the dispersion parameter. If `NULL`, it is inferred from object.
`REML`	optional logical value, passed to the `logLik` method indicating whether the restricted log-likelihood or log-likelihood should be used. The default is to use the method used for model estimation.

Details

Mallows' *Cp* statistic is the residual deviance plus twice the estimate of σ^2 times the residual degrees of freedom. It is closely related to AIC (and a multiple of it if the dispersion is known).

ICOMP (I for informational and COMP for complexity) penalizes the covariance complexity of the model rather than the number of free parameters directly.

$$\text{ICOMP} = -2\log L + k \, \log(\frac{tr(\Sigma)}{k}) - \log \mid \Sigma \mid$$

where log*L* is the log-likelihood, Σ is the variance-covariance matrix of parameters, *k* is the number of parameters, and *tr* is the trace of the matrix.

Value

If just one object is provided, returns a numeric value with the corresponding IC; if more than one object are provided, returns a `data.frame` with rows corresponding to the objects.

References

Mallows, C. L. (1973) Some Comments on CP. *Technometrics* 15: 661–675

Bozdogan, H. and Haughton, D.M.A. (1998) Information complexity criteria for regression models. *Comp. Stat. & Data Analysis* 28: 51-76

This implementation of ICOMP is based on functions in package **icomp** by J. Ferguson.

See Also

`AIC`, `AICc` and `BIC`

Miscellaneous	*Helper functions*

Description

`beta.weights` - computes standardized coefficients (beta weights) for a model;

`coeffs` - extracts model coefficients;

`getAllTerms` - extracts independent variable names from a model object;

`tTable` - extracts a table of coefficients, standard errors, and p-values from a model object;

`Weights` - calculates Akaike weights (normalized models likelihoods)

Usage

```
beta.weights(model)
coeffs(model)
getAllTerms(x, ...)
## S3 method for class 'terms'
getAllTerms(x, offset = TRUE, intercept = FALSE, ...)
tTable(model, ...)
Weights(aic, ...)

cbindDataFrameList(x)
rbindDataFrameList(x)
```

Arguments

model	a fitted model object
x	a fitted model object or a formula. for *bindDataFrameList, a list of data.frames
offset	should 'offset' terms be included?
intercept	should terms names include the intercept?
...	other arguments, often not used
aic	a vector of AIC (or other information criterion) values

Details

The functions `coeffs`, `getAllTerms` and `tTable` provide an interface between the model and `model.avg` (as well as `dredge`). Custom methods can be written to provide support for additional classes of models.

Note

`coeffs`'s value is in most cases identical to that returned by coef, the only difference is that it returns fixed effects' coefficients for mixed models.

Whimsically, the functions `*bindDataFrameList` are not exported from the name space, use `MuMIn:::cbindDataFrameList` to access them.

Author(s)

Kamil Bartoń

See Also

Vignette 'Extending **MuMIn**'s functionality' has information on how to use model selection functions with other model types.

```
mod.sel                          model selection table
```

Description

Builds a model selection table

Usage

```
mod.sel(object, ...)
model.sel(object, ...)

## S3 method for class 'model.selection'
mod.sel(object, rank = NULL, rank.args = NULL, ...)
## Default S3 method:
mod.sel(object, ..., rank = NULL, rank.args = NULL)
```

Arguments

object	A fitted model object, a list of such objects, or a `"model.selection"` object.
...	More fitted model objects
rank	Optional, custom rank function (information criterion) to use instead of `AICc`, e.g. `QAIC` or `BIC`, may be omitted if `object` is a model list returned by `get.models`.
rank.args	Optional `list` of arguments for the `rank` function. If one is an expression, an `x` within it is substituted with a current model.

Value

An object of class `"model.selection"` with columns containing useful information about each model: the coefficients, value of the information criterion used, Delta(IC) and weight.

Author(s)

Kamil Bartoń

See Also

dredge

Examples

```
data(Cement)
Cement$X1 <- cut(Cement$X1, 3)
Cement$X2 <- cut(Cement$X2, 2)

fm1 <- glm(formula = y ~ X + X1 + X2 * X3, data = Cement)
fm2 <- update(fm1, . ~ . - X - X1)
```

```
fm3 <- update(fm1, . ~ . - X2 - X3)

# ranked with AICc by default
mod.sel(fm1, fm2, fm3)


# ranked with BIC
mod.sel(fm1, fm2, fm3, rank=AIC, rank.args=alist(k=log(nobs(x))))
```

---

model.avg   *Model averaging*

Description

Model averaging based on an information criterion.

Usage

```
model.avg(object, ..., beta = FALSE, method = c("0", "NA"),
method.var = c("NA", "0"), rank = NULL, rank.args = NULL,
revised.var = TRUE)
```

Arguments

object   A fitted model object or a list of such objects. Alternatively an object of class
                model.selection. See 'Details'.

...   more fitted model objects

beta   Logical, should standardized coefficients be returned?

method, method.var
                The method of averaging estimators and their variance for parameters that are
                not common for all the models. Either "0" (default) or "NA". See 'Details'.

rank   Optional, custom rank function (information criterion) to use instead of AICc,
                e.g. QAIC or BIC, may be omitted if object is a model list returned by
                get.models or a model.selection object. See 'Details'.

rank.args   Optional list of arguments for the rank function. If one is an expression, an
                x within it is substituted with a current model.

revised.var   Logical, indicating whether to use revised formula for standard errors. See
                par.avg.

Details

model.avg has been tested to work with the following model classes:

- lm, glm
- gam (**mgcv**)
- lme, gls (**nlme**)
- lmer (**lme4**)

- `rlm`, `glm.nb` `polr` (**MASS**)
- `multinom` (**nnet**)
- `sarlm`, `spautolm` (**spdep**)
- `glmmML` (**glmmML**)
- `coxph` (**survival**)
- `unmarkedFit` (**unmarked**)

`model.avg` may be used with a list of models, or an object returned by `dredge`. In the latter case, the models from the model selection table are evaluated (with a call to `get.models`) prior to averaging. A warning is given if the subset argument is not provided, and the default `delta <= 4` will be used.

Other model types are also likely to be supported, in particular those inheriting from one of the above classes. See 'Details' section of the 'Miscellaneous' page to see how to provide support for other types of models.

With `method = "0"` (default) all predictors are averaged as if they were present in all models in the set, and the value of parameter estimate is taken to be 0 if it is not present in a particular model. If `method = "NA"`, the predictors are averaged only over the models in which they appear, which biases them away from zero. Analogically, the argument `method.var` defines the method of calculating unconditional variance estimators (in this case the default is `"NA"`).

Example in `link{par.avg}` shows comparison of the two methods for both coefficients and variance.

`rank` is found by a call to `match.fun` and typically is specified as a function or a symbol (e.g. a back-quoted name) or a character string specifying a function to be searched for from the environment of the call to lapply. `rank` must be a function able to accept model as a first argument and must always return a scalar.

Some generic methods such as `predict.averaging`, `coef`, `formula`, `residuals` and `vcov` are supported.

`logLik` method returns a list of `logLik` objects for the component models.

Value

An object of class `averaging` with following elements:

summary	a `data.frame` with deviance, AICc, Delta and weights for the component models.
coefficients, variance	
	matrices of component models' coefficients and their variances
variable.codes	
	names of the variables with numerical codes used in the summary
avg.model	the averaged model summary, (`data.frame` containing averaged coefficients, unconditional standard error, adjusted SE, and confidence intervals)
importance	the relative importance of the predictor variables: calculated as a sum of the Akaike weights over all of the models in which the parameter of interest appears.
beta	(logical) were standardized coefficients used?
term.names	character vector giving names of all terms in the model
residuals	the residuals (response minus fitted values).
x, formula	the model matrix and formula analogical to those that would be used in a single model.
method	how the missing terms were handled ("NA" or "0").
call	the matched call.

Note

From version 1.0.1, `print` method provides only a concise output (similarly as for `lm`). To print a full summary of the results use `summary` function. Confidence intervals can be obtained with `confint`.

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

See Also

See `par.avg` for details of averaged model calculation.

`dredge`, `get.models`. `QAIC` has examples of using custom rank function and prediction with confidence intervals.
`AICc` has examples of averaging models fitted by REML.

`modavg` in package **AICcmodavg**, and `coef.glmulti` in package **glmulti** also perform model averaging.

Examples

```
# Example from Burnham and Anderson (2002), page 100:
data(Cement)
lm1 <- lm(y ~ ., data = Cement)
dd <- dredge(lm1)
dd

#models with delta.aicc < 4
model.avg(get.models(dd, subset = delta < 4)) # get averaged coefficients

#or as a 95% confidence set:
top.models <- get.models(dd, cumsum(weight) <= .95)

model.avg(top.models) # get averaged coefficients
## Not run:
# The same result
model.avg(dd, cumsum(weight) <= .95)

## End(Not run)

## Not run:
# using BIC (Schwarz's Bayesian criterion) to rank the models
BIC <- function(x) AIC(x, k=log(length(residuals(x))))
mav <- model.avg(top.models, rank=BIC)

## End(Not run)
```

---

par.avg   *Parameter averaging*

Description

Averages single model coefficient based on provided weights

Usage

```
par.avg(x, se, weight, df = NULL, level = 1 - alpha, alpha = 0.05,
    revised.var = TRUE)
```

Arguments

`x`	vector of parameters
`se`	vector of standard errors
`weight`	vector of weights
`df`	(optional) vector of degrees of freedom
`alpha, level`	significance level for calculating confidence intervals
`revised.var`	logical, should the revised formula for standard errors be used? See 'Details'.

Details

Unconditional standard errors are square root of the variance estimator, calculated either according to the original formula in Burnham and Anderson (2002, p. 160, equation 4.7), or a newer, revised formula from Burnham and Anderson (2004, equation 4) (if `revised.var = TRUE`, this is the default). If degrees of freedom are given, the confidence intervals are based on adjusted standard error estimator (Burnham and Anderson 2002, page 164).

Value

`par.avg` returns a vector with named elements:

`Coefficient`	model coefficients
`SE`	unconditional standard error
`Adjusted SE`	adjusted standard error
`Lower CI, Upper CI`	
	unconditional confidence intervals

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

Burnham, K. P. and Anderson, D. R. (2004). *Multimodel inference - understanding AIC and BIC in model selection*. Sociological Methods & Research 33(2): 261-304.

See Also

model.avg for model averaging.

Examples

```
# Compare estimates and variance, model averaged with different methods
# ("NA", "0", see ?model.avg)

data(Cement)
fm1 <- lm(y ~ ., data = Cement)

avgmod <- model.avg(dredge(fm1), cumsum(weight) <= .95, method="NA")
weight <- avgmod$summary$Weight

na2z <- function(x) ifelse(is.na(x), 0, x) #helper function (changes NA to 0)

ret <- sapply(colnames(avgmod$coefficients), function(i) {
x <- avgmod$coefficients[, i]
se <- avgmod$se[, i]
df <- avgmod$dfs[, i]
rbind(
'coef(x)se(x)' = par.avg(x, se, weight, df),
'coef(0)se(x)' = par.avg(na2z(x), se, weight, df),
'coef(0)se(0)' = par.avg(na2z(x), na2z(se), weight, df)
)
}, simplify=FALSE)


if(require(graphics)) {

# Plot the comparison: coefficients +- confidence intervals
y <- 1:3
cap <- .1
op <- par(mfrow=c(2,3), mar=c(3,5,1,3), oma=c(3,0,3,0))
lapply(names(ret), function(i) {
x <- ret[[i]]
plot(y, x[,1], xlim=c(.5,3.5), ylim=range(x[,c(1,4,5)]), pch=20, cex=2,
ylab=i, xlab=NA, axes=FALSE, col=c(1,2,2))
segments(y, x[,4], y, x[,5], lty=2, col=c(1,1,2))
segments(y - cap, x[,4], y + cap, x[,4], col=c(1,1,2))
segments(y - cap, x[,5], y + cap, x[,5], col=c(1,1,2))
box(); axis(2); axis(1, at=y, labels=rownames(x))
abline(h=0, lty=3)
})
title("Comparison of model averaging methods",
sub="Cement hardening example", outer=TRUE, line=1)
par(op)

}
```

---

predict.averaging   *Predict Method for the Averaged Model*

Description

Model-averaged predictions with optional standard errors.

Usage

```
## S3 method for class 'averaging'
predict(object, newdata, se.fit = NULL, interval = NULL,
    type=c("link", "response"), ...)
```

Arguments

object	An object returned by model.avg.
newdata	An optional data frame in which to look for variables with which to predict. If omitted, the fitted values are used.
se.fit	logical, indicates if standard errors should be returned. This has any effect only if the predict methods for each of the component models support it.
interval	Currently not used.
type	Predictions on response scale are only possible if all component models use the same family. See predict.glm
...	Arguments to be passed to respective predict method (e.g. level for lme model).

Details

predict.averaging supports method = "NA" only for linear, fixed effect models. In other cases (e.g. nonlinear or mixed models), prediction is obtained using "brute force", i.e. by calling predict on each component model and weighted averaging the results, which is equivalent to assuming that missing coefficients equal zero (method = "0").

Predictions from generalized models are calculated by averaging predictions of each model on the link scale, followed by inverse transformation.

Besides predict and coef, other generic methods such as formula, residuals and vcov are supported.
logLik method returns a list of logLik objects for the component models.

Value

An object of class averaging with following elements:

summary	a data.frame with deviance, AICc, Delta and weights for the component models.
coefficients, variance	
	matrices of component models' coefficients and their variances
variable.codes	
	names of the variables with numerical codes used in the summary

avg.model	the averaged model summary, (data.frame containing averaged coefficients, unconditional standard error, adjusted SE, and confidence intervals)
importance	the relative importance of variables
beta	(logical) were standardized coefficients used?
term.names	character vector giving names of all terms in the model
residuals	the residuals (response minus fitted values).
x, formula	the model matrix and formula analogical to those that would be used in a single model.
method	how the missing terms were handled ("NA" or "0").
call	the matched call.

Note

predict.averaging relies on availability of the predict methods for the component model classes (except for lm/glm).

Author(s)

Kamil Bartoń

See Also

model.avg See par.avg for details of model-averaged parameter calculation.

Examples

```
require(graphics)

# Example from Burnham and Anderson (2002), page 100:
data(Cement)
lm1 <- lm(y ~ ., data = Cement)
dd <- dredge(lm1)
top.models <- get.models(dd, subset=cumsum(weight) <= .95)
avgm <- model.avg(top.models)

# helper function
nseq <- function(x, len=length(x)) seq(min(x, na.rm=TRUE),
    max(x, na.rm=TRUE), length=len)

# New predictors: X1 along the range of original data, other
# variables held constant at their means
newdata <- as.data.frame(lapply(lapply(Cement[1:5], mean), rep, 25))
newdata$X1 <- nseq(Cement$X1, nrow(newdata))

# Predictions from each of the models in a set:
pred <- sapply(top.models, predict, newdata=newdata)
# Add predictions from the models averaged using two methods:
pred <- cbind(pred,
    averaged.0=predict(avgm, newdata),
    averaged.NA=predict(update(avgm, method="NA"), newdata))

matplot(x=newdata$X1, y=pred, type="l", lwd=c(rep(1,ncol(pred)-2), 2, 2),
    xlab="X1", ylab="y")
```

```
legend("topleft",
    legend=c(lapply(top.models, formula),
        paste("Averaged model (method=", c("0", "NA"), ")", sep="")),
    col=1:6, lty=1:5, lwd=c(rep(1,ncol(pred)-2), 2, 2), cex = .75)


## Not run:
# Example with gam models (based on "example(gam)")
require(mgcv)
dat <- gamSim(1, n = 500, dist="poisson", scale=0.1)

gam1 <- gam(y ~ s(x0) + s(x1) + s(x2) +  s(x3) + (x1 + x2 + x3)^2,
    family = poisson, data = dat, method = "REML")

cat(dQuote(getAllTerms(gam1)), "\n")

# include only models with either smooth OR linear term (but not both)
# for each predictor variable:
dd <- dredge(gam1, subset=xor(`s(x1)`, x1) & xor(`s(x2)`, x2) &
    xor(`s(x3)`, x3))
# ...this may take a while.

subset(dd, cumsum(weight) < .95)

top.models <- get.models(dd, cumsum(weight) <= .95)

newdata <- as.data.frame(lapply(lapply(dat, mean), rep, 50))
newdata$x1 <- nseq(dat$x1, nrow(newdata))
pred <- cbind(
    sapply(top.models, predict, newdata=newdata),
    averaged=predict(model.avg(top.models), newdata))

matplot(x=newdata$x1, y=pred, type="l", xlab="x1", ylab="y"
    lwd=c(rep(1, ncol(pred) - 2), 2, 2))

## End(Not run)
```

QAIC	*Quasi AIC(c))*

Description

Calculates "quasi AIC" (or "quasi AICc") for one or several fitted model objects. This function is provided mainly as an example of custom rank function for use with `model.avg` and `dredge`

Usage

```
QAIC(object, ..., chat)
QAICc(object, ..., chat)
```

Arguments

`object`	a fitted model object.
`...`	optionally more fitted model objects.
`chat`	c - hat

Value

If just one object is provided, returns a numeric value with the corresponding QAIC; if more than one object are provided, returns a data.frame with rows corresponding to the objects.

Note

This implementation of QAIC uses model `deviance` rather than a likelihood itself. While this allows calculation also with models fitted using quasi-likelihood (where `logLik = NA`), the absolute values returned may differ from those obtained with the use of plain log-likelihood, since deviance is sometimes adjusted by a constant, so that the saturated model has deviance zero (see `glm`).

`dredge` will use `QAICc` instead of default `AICc` with `glm` with quasi* family.

Author(s)

Kamil Bartoń

See Also

`AICc`
`quasi` family used for models with over-dispersion.
`AIC` and `BIC` may also be used as a custom rank function in `dredge` and `model.avg`.

'Dealing with quasi- models in R', a vignette in the **bbmle** package.

Examples

```
# Based on "example(predict.glm)"

budworm <- data.frame(
    ldose = rep(0:5, 2),
    numdead = c(1, 4, 9, 13, 18, 20, 0, 2, 6, 10, 12, 16),
    sex = factor(rep(c("M", "F"), c(6, 6))))
budworm$SF = cbind(
    numdead = budworm$numdead,
    numalive = 20 - budworm$numdead)

budworm.lg <- glm(SF ~ sex*ldose, data = budworm, family = quasibinomial)

dd <- dredge(budworm.lg, rank = "QAIC",
    chat = summary(budworm.lg)$dispersion)
# Average all models
budworm.avg <- model.avg(get.models(dd, seq(nrow(dd))), method="NA")
#model.avg(mod[[1]], mod[[2]], rank = "QAIC", rank.args = list(chat = 1))

if(require(graphics)) {
```

```
plot(c(1,32), c(0,1), type = "n", xlab = "dose",
     ylab = "prob", log = "x")
text(2^budworm$ldose, budworm$numdead/20, as.character(budworm$sex))
ld <- seq(0, 5, 0.1)

newdata <- data.frame(ldose=ld, sex=factor(rep("M", length(ld)),
    levels=levels(budworm$sex)))

# Predictions from global model / Males
pred.lg <- predict(budworm.lg, newdata, se.fit=TRUE, type="response")
matplot(2^ld, cbind(pred.lg$fit, pred.lg$fit - (2 * pred.lg$se.fit),
    pred.lg$fit + (2 * pred.lg$se.fit)), add=TRUE, type="l", col=1)

# Predictions from averaged model / Males
pred.avg <- predict(budworm.avg, newdata, se.fit=TRUE, type="response")
matplot(2^ld, cbind(pred.avg$fit, pred.avg$fit - (2 * pred.avg$se.fit),
    pred.avg$fit + (2 * pred.avg$se.fit)), add=TRUE, type="l", col=2)

newdata$sex[] <- "F"

# Predictions from global model / Females
pred.lg <- predict(budworm.lg, newdata, se.fit=TRUE, type="response")
matplot(2^ld, cbind(pred.lg$fit, pred.lg$fit - (2 * pred.lg$se.fit),
    pred.lg$fit + (2 * pred.lg$se.fit)), add=TRUE, type="l", col=1)

# Predictions from averaged model / Females
pred.avg <- predict(budworm.avg, newdata, se.fit=TRUE, type="response")
matplot(2^ld, cbind(pred.avg$fit, pred.avg$fit - (2 * pred.avg$se.fit),
    pred.avg$fit + (2 * pred.avg$se.fit)), add=TRUE, type="l", col=2)

legend("bottomright", legend=c("full", "averaged"), title="Model",
       col=1:2, lty=1)

}
```

```
subset.model.selection
```
Subsetting model selection table

Description

Return subsets of a model selection table returned by `dredge`.

Usage

```
## S3 method for class 'model.selection'
subset(x, subset, select, recalc.weights = TRUE, ...)
## S3 method for class 'model.selection'
x[i, j, recalc.weights = TRUE, ...]
```

Arguments

x   a `model.selection` object to be subsetted.

subset,select

                   logical expressions indicating columns and rows to keep. See `subset`.

i,j   indices specifying elements to extract.

recalc.weights

                   logical value specyfying whether Akaike weights should be normalized across
                   the new set of models to sum to one.

...   further arguments passed to `[.data.frame`.

Value

A `model.selection` object containing only the selected models (rows). When columns are
selected (arguments `select` or `j` are provided), a plain `data.frame` is returned.

Note

Unlike the method for `data.frame`, extracting with only one index (i.e. `x[i]`) will select rows
rather than columns.

Author(s)

Kamil Bartoń

See Also

`dredge`, `subset` and `[.data.frame` for subsetting and extracting from data.frames.

Index