Package 'MuMIn'

November 5, 2011

Type Package

Title Multi-model inference

Version 1.6.0

Date 2011-11-07

Encoding UTF-8

Author Kamil Bartoń

Maintainer Kamil Bartoń <kamil.barton@go2.pl>

Description Model selection and model averaging based on information criteria (AICc and alike).

License GPL-2

Depends R (>= 2.12.0)

Imports stats

Suggests stats4, nlme, mgcv (>= 1.7.5), lme4 (>= 0.999375-
16), gamm4,MASS, nnet, spdep, survival, unmarked, glmmML

LazyLoad yes

BuildVignettes true

R topics documented:

MuMIn-package   *Multi-model inference*

Description

The package MuMIn contains functions to streamline model selection and perform model averaging based on information criteria (AIC and alike).

Details

User-level functions include:

dredge performs automated model selection with subsets of the supplied 'global' model, and op-tional choices of other model properties (such as different link functions). The set of models may be generated either with 'all possible' combinations, or tailored according to the conditions specified.

model.sel creates a model selection table from hand-picked models.

model.avg calculates model averaged parameters, with standard errors and confidence intervals.

AICc calculates second-order Akaike information criterion.

Model selection can be done according to any information criterion, such as AIC, AICc, BIC, QAIC, ICOMP or Mallows' Cp. Most of R's common model object classes are supported.

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed. New York, Springer-Verlag.

See Also

AIC, step or stepAIC for stepwise model selection by AIC.

Examples

```
data(Cement)

fm1 <- lm(y ~ ., data = Cement)

ms1 <- dredge(fm1)
confset.d4 <- get.models(ms1, subset = delta < 4)
```

```
model.avg(confset.d4)

confset.95p <- get.models(ms1, cumsum(weight) <= .95)
avgmod.95p <- model.avg(confset.95p)
summary(avgmod.95p)
confint(avgmod.95p)
```

AICc   *Second-order Akaike Information Criterion*

Description

Calculate second-order Akaike information criterion for one or several fitted model objects (AIC_c, AIC for small samples).

Usage

```
AICc(object, ..., k = 2, REML = NULL)
```

Arguments

object   a fitted model object for which there exists a `logLik` method, or a `logLik` object.

...   optionally more fitted model objects.

k   the 'penalty' per parameter to be used; the default k = 2 is the classical AIC.

REML   optional logical value, passed to the `logLik` method indicating whether the restricted log-likelihood or log-likelihood should be used. The default is to use the method used for model estimation.

Value

If just one object is provided, returns a numeric value with the corresponding AIC_c; if more than one object are provided, returns a `data.frame` with rows corresponding to the objects and columns representing the number of parameters in the model (df) and AIC_c.

Note

AIC_c should be used instead AIC when the the sample size is small in comparison to the number of estimated parameters (Burnham & Anderson 2002 recommend it's use when $n/K < 40$).

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed. New York, Springer-Verlag.

Hurvich, C. M. and Tsai, C.-L. (1989) Regression and time series model selection in small samples, *Biometrika* 76: 297–307.

See Also

Akaike's An Information Criterion: AIC

Other implementations: AICc in package **AICcmodavg**, AICc in package **bbmle** and aicc in package **glmulti**

Examples

```
#Model-averaging mixed models

library(nlme)
data(Orthodont, package = "nlme")

# Fit model by REML
fm2 <- lme(distance ~ Sex*age, data = Orthodont,
    random = ~ 1|Subject / Sex, method = "REML")

# Model selection: ranking by AICc using ML
ms2 <- dredge(fm2, trace = TRUE, rank = "AICc", REML = FALSE)

(attr(ms2, "rank.call"))

# Get the models (fitted by REML, as in the global model)
fmList <- get.models(ms2, 1:4)

# Because the models originate from 'dredge(..., rank=AICc, REML=FALSE)',
# the default weights in 'model.avg' are ML based:
model.avg(fmList)

# same result
#model.avg(fmList, rank = "AICc", rank.args = list(REML=FALSE))
```

Beetle   *Flour beetle mortality data*

Description

Mortality of flour beetles (*Tribolium confusu*) due to exposure to gaseous carbon disulfide CS_2, from Bliss (1935).

Usage

```
data(Beetle)
```

Format

Beetle is a data frame with 5 elements.

dose The dose of CS_2 in mg/L
n.tested Number of beetles tested
n.killed Number of beetles killed
Prop A matrix with two columns named **n.killed** and **n.survived**
mortality Observed mortality rate.

Source

Bliss C. I. (1935) The calculation of the dosage-mortality curve. *Annals of Applied Biology*, 22: 134–167.

References

Burnham, K. P. and Anderson, D. R. (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed. New York, Springer-Verlag.

Examples

```
# "Logistic regression example"
# from Burnham & Anderson (2002) chapter 4.11

data(Beetle)

# Fit a global model with all the considered variables
globmod <- glm(Prop ~ dose + I(dose^2) + log(dose) + I(log(dose)^2),
    data = Beetle, family = binomial)


# A logical expression defining the subset of models to use:
# * either log(dose) or dose
# * the quadratic terms can appear only together with linear terms
msubset <- expression(xor(dose, `log(dose)`) & (dose | !`I(dose^2)`)
    & (`log(dose)` | !`I(log(dose)^2)`))


# Table 4.6

# Use 'varying' argument to fit models with different link functions
# Note the use of 'alist' rather than 'list' in order to keep the 'family'
# objects unevaluated
varying.link <- list(family = alist(
    logit = binomial("logit"),
    probit = binomial("probit"),
    cloglog = binomial("cloglog")
    ))

(ms12 <- dredge(globmod, subset = msubset, varying = varying.link, rank = AIC))

# Table 4.7 "models justifiable a priori"
(ms3 <- subset(ms12, has(dose, !`I(dose^2)`)))

# The same result, but would fit the models again:
# ms3 <- update(ms12, update(globmod, . ~ dose), subset = , fixed = ~dose)

mod3 <- get.models(ms3, 1:3)

# Table 4.8. Predicted mortality probability at dose 40.

# calculate confidence intervals on logit scale
logit.ci <- function(p, se, quantile = 2) {
    C. <- exp(quantile * se / (p * (1 - p)))
    p /(p + (1 - p) * c(C., 1/C.))
```

```
}

pred <- sapply(mod3, predict, newdata = list(dose = 40), se.fit = TRUE,
    type = "response")
pred <- apply(pred, 1, unlist)[, 1:2] # simplify

# build the table
tab <- rbind(pred, par.avg(pred[, "fit"], pred[, "se.fit"], Weights(ms3),
    revised.var = FALSE)[1:2])
tab <- cbind(
    c(Weights(ms3), NA),
    tab,
    matrix(logit.ci(tab[,"fit"], tab[,"se.fit"], quantile = c(rep(1.96, 3), 2)),
    ncol = 2)
    )
colnames(tab) <- c("Akaike weight", "Predicted(40)", "SE", "Lower CI", "Upper CI")
rownames(tab) <- c(as.character(ms3$family), "model averaged")

print(tab, digits = 3, na.print = "")


# Figure 4.3
newdata <- list(dose = seq(min(Beetle$dose), max(Beetle$dose), length.out = 25))
matplot(newdata$dose, sapply(mod3, predict, newdata, type="response"),
    type = "l", xlab = quote(list("Dose of" ~ CS[2],(mg/L))),
    ylab = "Mortality", col = 2:4, lty = 3, lwd = 1
)

# add model-averaged prediction with CI, using the same method as above
pred <- lapply(mod3, predict, newdata, type = "response", se.fit = TRUE)
pred.y <- sapply(pred, "[[", "fit")
pred.se <- sapply(pred, "[[", "se.fit")
avpred <- sapply(1:25, function(i) par.avg(pred.y[i, ], pred.se[i, ],
    weight = Weights(ms3), revised.var = FALSE)[1:2])
avci <- matrix(logit.ci(avpred[1, ], avpred[2, ], quantile = 2), ncol = 2)
matplot(newdata$dose, cbind(avpred[1, ], avci), type = "l", add = TRUE, lwd = 1,
    lty = c(1, 2, 2), col = 1)
legend("topleft", NULL, c(as.character(ms3$family), expression('averaged' %+-% CI)),
    lty = c(3, 3, 3, 1), col = c(2:4, 1))
```

Cement	*Cement hardening data*

Description

Cement hardening data from Woods et al (1939).

Usage

```
data(Cement)
```

Format

Cement is a data frame with 5 variables. x1-x4 are four predictor variables expressed as a percentage of weight.

X1 calcium aluminate

X2 tricalcium silicate

X3 tetracalcium alumino ferrite

X4 dicalcium silicate

y calories of heat evolved per gram of cement after 180 days of hardening.

Source

Woods H., Steinour H.H., Starke H.R. (1932) Effect of composition of Portland cement on heat evolved during hardening. *Industrial & Engineering Chemistry* 24, 1207-1214.

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed. New York, Springer-Verlag.

---

dredge	*Automated model selection*

Description

Generate a set of models with combinations of the terms in the global model, with optional rules for inclusion.

Usage

```
dredge(global.model, beta = FALSE, evaluate = TRUE, rank = "AICc",
    fixed = NULL, m.max = NA, m.min = 0, subset, marg.ex = NULL,
    trace = FALSE, varying, extra, ...)


## S3 method for class 'model.selection'
print(x, abbrev.names = TRUE, ...)
```

Arguments

global.model	a fitted 'global' model object. See 'Details' for a list of supported types.
beta	logical, should standardized coefficients be returned?
evaluate	whether to evaluate and rank the models. If FALSE, a list of model calls is returned.
rank	optional custom rank function (information criterion) to be used instead AICc, e.g. QAIC or BIC. See 'Details'.
fixed	optional, either a single sided formula or a character vector giving names of terms to be included in all models.

m.max, m.min	optionally the maximum and minimum number of terms in a single model (excluding the intercept), m.max defaults to the number of terms in global.model.
subset	logical expression describing models to keep in the resulting set. See 'Details'.
marg.ex	a character vector specifying names of variables for which NOT to check for marginality restrictions when generating model formulas. If this argument is set to TRUE, all combinations of terms are used (i.e. no checking).
trace	if TRUE, all calls to the fitting function (i.e. updated global.model calls) are printed before actual fitting takes place.
varying	optionally, a named list describing the additional arguments to vary between the generated models. Names are the names of the arguments, and each item provides a list of choices. Complex items in the choice list (such as family objects) should be either named (uniquely) or quoted (unevaluated, e.g. using alist, see quote), otherwise it may produce rather unpleasant effects. See example in Beetle.
extra	optional additional statistics to include in the result, provided as functions, function names or a list of such (best if named or quoted). Similarly as in rank argument, each function must accept fitted object model as an argument and return (object coercible to) a numeric vector. These can be e.g. additional information criterions or goodness-of-fit statistics. The character string "R^2" is treated in a special way, and will add a likelihood-ratio based R^2 to the resulting table (this is more efficient than using r.squaredLR directly).
x	a model.selection object, returned by dredge.
abbrev.names	should variable names be abbreviated when printing? (useful with many variables).
...	optional arguments for the rank function. Any can be an expression (of mode call), in which case any x within it will be substituted with a current model.

Details

Fitted model objects that can be used as a global.model include those returned by lm, glm (package **stats**); gam, gamm (**mgcv**); gamm4 (**gamm4**); lme, gls (**nlme**); lmer (**lme4**); rlm, glm.nb, polr (**MASS**); multinom (**nnet**); sarlm, spautolm (**spdep**); glmmML (**glmmML**); coxph, survreg (**survival**); rq (**quantreg**); and most models from package **unmarked**. gamm and gamm4 should be evaluated *via* the wrapper MuMIn::gamm.

Models are run one by one by repeated evaluation of modified calls to global.model (in a similar fashion to that of update. This method, while robust in that it can be applied to a variety of different model objects is not very efficient, and may be time-intensive.

Note that the number of combinations grows exponentially with number of predictor variables (2^N). Because there can be potentially a large number of models to evaluate, to avoid memory overflow the fitted model objects are not stored in the result. To get (a subset of) the models, use get.models on the object returned by dredge.

Handling interactions, dredge respects marginality constraints, so "all possible combinations" do not include models containing interactions without their respective main effects. This behaviour can be altered by marg.ex argument. It can be used to allow for simple nested designs. For example, with global model of form a / (x + z), use marg.ex = "a" and fixed = "a".

rank is found by a call to match.fun and may be specified as a function or a symbol (e.g. a backquoted name) or a character string specifying a function to be searched for from the environment of the call to dredge. Function rank must be able to accept model as a first argument and must always return a scalar. Typical choice for rank would be "AIC", "QAIC" or "BIC" (**stats** or **nlme**).

The argument subset acts in a similar fashion to that in the function subset for data.frames: the model terms can be referred to by name as variables in the expression, with the difference that they are always logical (i.e. TRUE if a term exists in the model). The expression can contain any of the global.model terms (use getAllTerms(global.model) to list them). It can have a form of an unevaluated call, expression object, or a one sided formula. See 'Examples'. Compound model terms (such as 'as-is' expressions within I() or the smooths in gam) should be treated as non-syntactic names and enclosed in back-ticks (see Quotes). Mind the spacing, names must match exactly the term names in model's formula. To simply keep certain variables in all models, use of fixed is preferred.

Use of na.action = na.omit (R's default) in global.model should be avoided, as it results with sub-models fitted to different data sets, if there are missing values. In versions >= 0.13.17 a warning is given in such a case.

Value

dredge returns an object of class model.selection, being a data.frame with models' coefficients (or presence/NA for factors), *df* - number of parameters, log-likelihood, the information criterion value, delta-IC and *Akaike weight*. Models are ordered according to the used information criterion (lowest on top), specified by rank.

The attribute "calls" is a list containing the model calls used (arranged in the same order as the models). Other attributes: "global" - the global.model object, "rank" - the rank function used, and "call" - the matched call.

Note

Users should keep in mind the hazards that a "thoughtless approach" of evaluating all possible models poses. Although this procedure is in certain cases useful and justified, it may result in selecting a spurious "best" model, due to model selection bias.

"Let the computer find out" is a poor strategy and usually reflects the fact that the researcher did not bother to think clearly about the problem of interest and its scientific setting (Burnham and Anderson, 2002).

Author(s)

Kamil Bartoń

See Also

get.models, model.avg.

There are subset and plot methods.

Possible alternatives: glmulti in package **glmulti** and bestglm (**bestglm**), or aictab (**AICcmodavg**) and ICtab (**bbmle**) for "hand-picked" model selection tables.

regsubsets in package **leaps** also performs all-subsets regression.

Examples

```
# Example from Burnham and Anderson (2002), page 100:
data(Cement)
fm1 <- lm(y ~ ., data = Cement)
dd <- dredge(fm1)
subset(dd, delta < 4)
```

```
# Visualize the model selection table:
if(require(graphics))
plot(dd)


# Model average models with delta AICc < 4
model.avg(get.models(dd, subset = delta < 4))

#or as a 95% confidence set:
confset.95p <- get.models(dd, cumsum(weight) <= .95)

model.avg(confset.95p) # get averaged coefficients

#'Best' model
summary(confset.95p[[1]])

## Not run:
# Examples of using 'subset':
# exclude models containing both X1 and X2
dredge(fm1, subset = !(X1 & X2))
# keep only models containing X3
dredge(fm1, subset = ~ X3) # subset as a formula
dredge(fm1, subset = expression(X3)) # subset as expression object
# the same, but more effective:
dredge(fm1, fixed = "X3")

#Reduce the number of generated models, by including only those with
# up to 2 terms (and intercept)
dredge(fm1, m.max = 2)

## End(Not run)


# Add R^2 and F-statistics, use the 'extra' argument
dredge(fm1, m.max = 1, extra = c("R^2", F = function(x) summary(x)$fstatistic[[1]]))

# with summary statistics:
dredge(fm1, m.max = 1, extra = list(
    "R^2", "*" = function(x) {
        s <- summary(x)
        c(Rsq = s$r.squared, adjRsq = s$adj.r.squared, F = s$fstatistic[[1]])
    })
)


# with other information criterions:

# there is no BIC in R < 2.13.0, so need to add it:
if(!exists("BIC", mode="function"))
    BIC <- function(object, ...) AIC(object, k = log(length(resid(object))))

dredge(fm1, m.max = 1, extra = alist(AIC, BIC, ICOMP, Cp))
```

gamm-wrapper   *Updateable* gamm

Description

Enables updating of the model objects fitted by gamm and gamm4 from packages **mgcv** and **gamm4**.

Usage

```
gamm(formula, random = NULL, ..., lme4 = inherits(random, "formula"))
```

Arguments

formula, random, ...

> arguments passed to gamm or gamm4.

lme4   logical, if TRUE gamm4 is used rather than gamm. If TRUE, the random argument
                    must be provided as a formula.

Details

This function is just a wrapper for gamm and gamm4. The only purpose of it is to add a call compo-
nent, that is not provided by gamm* as such. It allows update on the returned object, so also makes
possible using it in model selection with dredge.

This is only a temporary workaround and it is likely be removed soon.

Value

Depending on the value of the 'lme4' switch, either a gamm or gamm4 fitted model object. The only
difference from the original object is an addition of the call component.

Note

To assure gamm is called *via* this wrapper in case it is masked by the original gamm from **mgcv** (when
MuMIn was loaded after **mgcv**), use MuMIn::gamm.

Author(s)

Kamil Bartoń

See Also

gamm and gamm4

---

get.models   *Get models*

Description

Generate a list of fitted model objects from a `model.selection` table.

Usage

```
get.models(object, subset = delta <= 4, ...)
```

Arguments

object	object returned by [dredge](#).
subset	subset of models, an expression evaluated within the model selection table, see [subset method](#).
...	additional parameters passed to `update`. For example, in lme/lmer one may want to use `method = "REML"` while using "ML" for model selection.

Value

[list](#) of fitted model objects

Author(s)

Kamil Bartoń

See Also

[dredge](#), [model.avg](#)

Examples

```
# Mixed models:

require(nlme)
fm2 <- lme(distance ~ age + Sex, data = Orthodont,
    random = ~ 1 | Subject, method="ML")
ms2 <- dredge(fm2)

# Get top-most models, but fitted by REML:
(confset.d4 <- get.models(ms2, subset = delta < 4, method = "REML"))
```

---

importance   *Relative variable importance*

Description

Sum of 'Akaike weights' over all models including the explanatory variable.

Usage

```
importance(x)
```

Arguments

x   Either a list of fitted model objects, or a "model.selection" or "averaging" object.

Value

a numeric vector of relative importance values, named as the predictor variables.

Author(s)

Kamil Bartoń

See Also

[Weights](#)

[dredge](#), [model.avg](#), [mod.sel](#)

Examples

```
# Generate some models
data(Cement)
fm1 <- lm(y ~ ., data = Cement)
ms1 <- dredge(fm1)

# Importance can be calculated/extracted from various objects:
importance(ms1)
## Not run:
importance(subset(mod.sel(ms1), delta <= 4))
importance(model.avg(ms1, subset = delta <= 4))
importance(subset(ms1, delta <= 4))
importance(get.models(ms1, delta <= 4))

## End(Not run)

# Re-evaluate the importances according to BIC
# note that re-ranking involves fitting the models again

# 'nobs' is not used here for backwards compatibility
lognobs <- log(length(resid(fm1)))
```

```
importance(subset(mod.sel(ms1, rank = AIC, rank.args = list(k = lognobs)),
cumsum(weight) <= .95))

# This gives a different result than previous command, because 'subset' is
# applied to the original selection table that is ranked with 'AICc'
importance(model.avg(ms1, rank = AIC, rank.args = list(k = lognobs),
subset = cumsum(weight) <= .95))
```

Information criteria   *Various information criteria*

Description

Calculate Mallows' *Cp* and Bozdogan's ICOMP information criterion.

Usage

```
Cp(object, dispersion = NULL)
ICOMP(object, ..., REML = NULL)
```

Arguments

object	a fitted model object (in case of ICOMP, a logLik method must exist for the object).
...	optionally more fitted model objects.
dispersion	the dispersion parameter. If NULL, it is inferred from object.
REML	optional logical value, passed to the logLik method indicating whether the restricted log-likelihood or log-likelihood should be used. The default is to use the method used for model estimation.

Details

Mallows' *Cp* statistic is the residual deviance plus twice the estimate of σ^2 times the residual degrees of freedom. It is closely related to AIC (and a multiple of it if the dispersion is known).

ICOMP (I for informational and COMP for complexity) penalizes the covariance complexity of the model, rather than the number of parameters directly.

Value

If just one object is provided, the functions return a numeric value with the corresponding IC; otherwise a data.frame with rows corresponding to the objects is returned.

References

Mallows, C. L. (1973) Some Comments on CP. *Technometrics* 15: 661–675.

Bozdogan, H. and Haughton, D.M.A. (1998) Information complexity criteria for regression models. *Comp. Stat. & Data Analysis* 28: 51-76.

See Also

AIC and BIC in **stats**, AICc

Miscellaneous   *Helper functions*

Description

beta.weights - computes standardized coefficients (beta weights) for a model;

coeffs - extracts model coefficients;

getAllTerms - extracts independent variable names from a model object;

tTable - extracts a table of coefficients, standard errors, and p-values from a model object.

Usage

```
beta.weights(model)
coeffs(model)
getAllTerms(x, ...)
## S3 method for class 'terms'
getAllTerms(x, offset = TRUE, intercept = FALSE, ...)
tTable(model, ...)

cbindDataFrameList(x)
rbindDataFrameList(x)
```

Arguments

model	a fitted model object.
x	a fitted model object or a [formula](#). for *bindDataFrameList, a list of data.frames.
offset	should 'offset' terms be included?
intercept	should terms names include the intercept?
...	other arguments, often not used.

Details

The functions coeffs, getAllTerms and tTable provide interface between the fitted model object and model.avg (and dredge). Custom methods can be written to provide support for additional classes of models. The vignette 'Extending **MuMIn**'s functionality' describes it in more detail.

Note

coeffs's value is in most cases identical to that returned by [coef](#), the only difference being it returns fixed effects' coefficients for mixed models.

Whimsically, the functions *bindDataFrameList are not exported from the name space, use MuMIn:::cbindDataFrameList to access them.

Author(s)

Kamil Bartoń

---

model.avg   *Model averaging*

Description

Model averaging based on an information criterion.

Usage

```
model.avg(object, ..., beta = FALSE, rank = NULL, rank.args = NULL,
    revised.var = TRUE)
```

Arguments

object   A fitted model object or a list of such objects. Alternatively an object of class
             model.selection. See 'Details'.

...   more fitted model objects.

beta   Logical, should standardized coefficients be returned?

rank   Optional, custom rank function (information criterion) to use instead of AICc,
             e.g. BIC or QAIC, may be omitted if object is a model list returned by get.models
             or a model.selection object. See 'Details'.

rank.args   Optional list of arguments for the rank function. If one is an expression, an x
             within it is substituted with a current model.

revised.var Logical, indicating whether to use revised formula for standard errors. See
             par.avg.

Details

model.avg has been tested to work with the fitted objects from following modelling functions:
lm, glm; gam, gamm (**mgcv**); gamm4 (**gamm4**); lme, gls (**nlme**); lmer (**lme4**); rlm, glm.nb, polr
(**MASS**); multinom (**nnet**); sarlm, spautolm (**spdep**); glmmML (**glmmML**); coxph, survreg (**survival**); and several models within the class unmarkedFit (**unmarked**). Other classes are also likely
to be supported, in particular those inheriting from one of the above classes. See package vignette
'Extending **MuMIn**'s functionality' for a demonstration on how to provide support for other types
of models.

model.avg may be used with a list of models, but also directly a model.selection object returned
by dredge. In the latter case, the models from the model selection table are evaluated (with a call
to get.models) prior to averaging. A warning is given if the subset argument is not provided, and
the default delta <= 4 will be used.

rank is found by a call to match.fun and typically is specified as a function or a symbol (e.g. a back-
quoted name) or a character string specifying a function to be searched for from the environment of
the call to lapply. rank must be a function able to accept model as a first argument and must always
return a scalar.

Several standard methods for fitted model objects exist for class averaging, including summary,
predict, coef, confint, formula, residuals, vcov. The coef method a accepts argument full,
if set to TRUE the full model-averaged coefficients are returned, rather than subset-averaged ones.
logLik returns a list of logLik objects for the component models.

Value

An object of class `averaging` is a list with components:

summary	a `data.frame` with deviance, AICc, Delta and weights for the component models.
coefficients, se, dfs	matrices of component models' coefficients, their standard errors, and degrees of freedom.
coef.shrinkage	a vector of full model-averaged coefficients, see 'Note'.
variable.codes	names of the variables with numerical codes used in `summary`.
avg.model	the model averaged parameters. A `data.frame` containing averaged coefficients, unconditional standard error, adjusted SE (if *df*s are available) and z-values (coefficient and SE) and significance (assuming a normal error distribution).
importance	relative importance of the predictor variables, calculated as a sum of the Akaike weights over all of the models in which the parameter of interest appears.
term.names	character vector giving names of all terms in the model.
x, formula	the model matrix and formula corresponding to the one that would be used in a single model. `formula` contains only the averaged coefficients.
residuals	model averaged residuals (response minus fitted values).
call	the matched call.

In addition, the object has following attributes:

mList	a list of component model objects.
beta	logical, were standardized coefficients used?
revised.var	if TRUE, the standard errors were calculated with the revised formula (See `par.avg`).

Note

The 'subset' (or 'conditional') average only averages over the models where the parameter appears. An alternative, the 'full' average assumes that a variable is included in every model, but in some models the corresponding coefficient is set to zero. Unlike the 'subset average', it does not have a tendency of biasing the value away from zero. It is, however, an unresolved issue how the variance of this estimate should be calculated, therefore the standard errors and confidence interval are returned only for the subset-averaged coefficients (as from version >= 1.5.0 argument `method` is no longer accepted, use of it throws an error).

Averaging models with different contrasts for the same factor would yield nonsense results.

From version 1.0.1, `print` method provides only a concise output (similarly as for `lm`). To print a full summary of the results use `summary` function. Confidence intervals can be obtained with `confint`.

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed. New York, Springer-Verlag.

See Also

See par.avg for more details of model averaged parameter calculation.

dredge, get.models
AICc has examples of averaging models fitted by REML.

modavg in package **AICcmodavg**, and coef.glmulti in package **glmulti** also perform model averaging.

Examples

```
# Example from Burnham and Anderson (2002), page 100:
data(Cement)
fm1 <- lm(y ~ ., data = Cement)
(ms1 <- dredge(fm1))

#models with delta.aicc < 4
summary(model.avg(get.models(ms1, subset = delta < 4))) # get averaged coefficients

#or as a 95% confidence set:
confset.95p <- get.models(ms1, cumsum(weight) <= .95)

avgmod.95p <- model.avg(confset.95p) # get averaged coefficients
confint(avgmod.95p)


## Not run:
# The same result
model.avg(ms1, cumsum(weight) <= .95)

## End(Not run)

## Not run:
# using BIC (Schwarz's Bayesian criterion) to rank the models
BIC <- function(x) AIC(x, k=log(length(residuals(x))))
model.avg(confset.95p, rank=BIC)
# the same result, using AIC directly, with argument k
# 'x' in a quoted 'rank' argument is substituted with a model object
# (in this case it does not make much sense as the number of observations is
# common to all models)
model.avg(confset.95p, rank=AIC, rank.args=alist(k=log(length(residuals(x)))))

## End(Not run)
```

---

model.sel   *model selection table*

Description

Build a model selection table.

Usage

```
mod.sel(object, ...)
model.sel(object, ...)

## S3 method for class 'model.selection'
mod.sel(object, rank = NULL, rank.args = NULL, ...)
## Default S3 method:
mod.sel(object, ..., rank = NULL, rank.args = NULL)
```

Arguments

`object`	A fitted model object, a list of such objects, or a `"model.selection"` object.
`...`	More fitted model objects.
`rank`	Optional, custom rank function (information criterion) to use instead of AICc, e.g. `QAIC` or `BIC`, may be omitted if `object` is a model list returned by `get.models`.
`rank.args`	Optional `list` of arguments for the `rank` function. If one is an expression, an `x` within it is substituted with a current model.

Value

An object of class `"model.selection"` with columns containing useful information about each model: the coefficients, value of the information criterion used, Delta(IC) and weight.

Author(s)

Kamil Bartoń

See Also

dredge

Examples

```
data(Cement)
Cement$X1 <- cut(Cement$X1, 3)
Cement$X2 <- cut(Cement$X2, 2)

fm1 <- glm(formula = y ~ X1 + X2 * X3, data = Cement)
fm2 <- update(fm1, . ~ . - X1 - X2)
fm3 <- update(fm1, . ~ . - X2 - X3)

# ranked with AICc by default
mod.sel(fm1, fm2, fm3)


# ranked with BIC
mod.sel(fm1, fm2, fm3, rank=AIC, rank.args=alist(k=log(nobs(x))))
```

---

par.avg   *Parameter averaging*

Description

Averages a single model coefficient based on provided weights.

Usage

```
par.avg(x, se, weight, df = NULL, level = 1 - alpha, alpha = 0.05,
    revised.var = TRUE, adjusted = TRUE)
```

Arguments

x	vector of parameters.
se	vector of standard errors.
weight	vector of weights.
df	(optional) vector of degrees of freedom.
alpha, level	significance level for calculating confidence intervals.
revised.var	logical, should the revised formula for standard errors be used? See 'Details'.
adjusted	logical, should the inflated standard errors be calculated? See 'Details'.

Details

Unconditional standard errors are square root of the variance estimator, calculated either according to the original equation in Burnham and Anderson (2002, equation 4.7), or a newer, revised formula from Burnham and Anderson (2004, equation 4) (if revised.var = TRUE, this is the default). If adjusted = TRUE (the default) and degrees of freedom are given, the confidence intervals are based on adjusted standard error estimator and confidence intervals with improved coverage (see Burnham and Anderson 2002, section 4.3.3).

Value

par.avg returns a vector with named elements:

Coefficient	model coefficients,
SE	unconditional standard error,
Adjusted SE	adjusted standard error,
Lower CI, Upper CI	
	unconditional confidence intervals.

Author(s)

Kamil Bartoń

References

Burnham, K. P. and Anderson, D. R (2002) *Model selection and multimodel inference: a practical information-theoretic approach*. 2nd ed.

Burnham, K. P. and Anderson, D. R. (2004). *Multimodel inference - understanding AIC and BIC in model selection*. Sociological Methods & Research 33(2): 261-304.

See Also

model.avg for model averaging.

---

predict.averaging   *Predict method for the averaged model*

Description

Model-averaged predictions with optional standard errors.

Usage

```
## S3 method for class 'averaging'
predict(object, newdata = NULL, se.fit = FALSE,
    interval = NULL, type = c("link", "response"), full = TRUE, ...)
```

Arguments

object	An object returned by model.avg.
newdata	An optional data frame in which to look for variables with which to predict. If omitted, the fitted values are used.
se.fit	logical, indicates if standard errors should be returned. This has any effect only if the predict methods for each of the component models support it.
interval	Currently not used.
type	Predictions on response scale are only possible if all component models use the same family. See predict.glm.
full	If TRUE, the full model averaged coefficients are used (only if se.fit = FALSE and the component objects are a result of lm).
...	Arguments to be passed to respective predict method (e.g. level for lme model).

Details

If all the component models are oridinary linear models, the prediction can be made either with the full averaged coefficients (the argument full = TRUE this is the default) or subset-averaged coefficients. Otherwise the prediction is obtained by calling predict on each component model and weighted averaging the results, which corresponds to the assumption that all predictors are present in all models, but those not estimated are equal zero. See 'Note' in model.avg. Predictions from component models with standard errors are passed to par.avg and averaged in the same way as the coefficients.

Predictions on the response scale from generalized models are calculated by averaging predictions of each model on the link scale, followed by inverse transformation.

Value

If `se.fit = FALSE`, a vector of predictions. If `se = TRUE`, a list with components: `fit` containing the predictions, and `se.fit` with the estimated standard errors.

Note

This method relies on availability of the `predict` methods for the component model classes (except when all component models are of class `lm`).

Author(s)

Kamil Bartoń

See Also

[model.avg](#) See [par.avg](#) for details of model-averaged parameter calculation.

Examples

```
require(graphics)

# Example from Burnham and Anderson (2002), page 100:
data(Cement)
fm1 <- lm(y ~ X1 + X2 + X3 + X4, data = Cement)

ms1 <- dredge(fm1)
confset.95p <- get.models(ms1, subset=cumsum(weight) <= .95)
avgm <- model.avg(confset.95p)

nseq <- function(x, len = length(x)) seq(min(x, na.rm = TRUE),
    max(x, na.rm=TRUE), length = len)

# New predictors: X1 along the range of original data, other
# variables held constant at their means
newdata <- as.data.frame(lapply(lapply(Cement[1:4], mean), rep, 25))
newdata$X1 <- nseq(Cement$X1, nrow(newdata))

n <- length(confset.95p)

# Predictions from each of the models in a set, and with averaged coefficients
pred <- data.frame(
model = sapply(confset.95p, predict, newdata = newdata),
averaged.subset = predict(avgm, newdata, full = FALSE),
    averaged.full = predict(avgm, newdata, full = TRUE)
)

opal <- palette(c(topo.colors(n), "black", "red", "orange"))
matplot(newdata$X1, pred, type = "l",
lwd = c(rep(2,n),3,3), lty = 1,
    xlab = "X1", ylab = "y", col=1:7)

# For comparison, prediction obtained by averaging predictions of the component
# models
pred.se <- predict(avgm, newdata, se.fit = TRUE)
y <- pred.se$fit
```

```
ci <- pred.se$se.fit  * 2
matplot(newdata$X1, cbind(y, y - ci, y + ci), add = TRUE, type="l",
lty = 2, col = n + 3, lwd = 3)

legend("topleft",
    legend=c(lapply(confset.95p, formula),
        paste(c("subset", "full"), "averaged"), "averaged predictions + CI"),
    lty = 1, lwd = c(rep(2,n),3,3,3),  cex = .75, col=1:8)

palette(opal)
```

QAIC   *Quasi AIC or AICc*

Description

Calculate a modification of Akaike's Information Criterion for overdispersed count data (or its version corrected for small sample, "quasi-AIC$_c$"), for one or several fitted model objects.

Usage

```
QAIC(object, ..., chat, k = 2)
QAICc(object, ..., chat, k = 2)
```

Arguments

object	a fitted model object.
...	optionally, more fitted model objects.
chat	$\hat{c}$, the variance inflation factor.
k	the 'penalty' per parameter.

Value

If only one object is provided, returns a numeric value with the corresponding QAIC or QAIC$_c$; otherwise returns a data.frame with rows corresponding to the objects.

Note

$\hat{c}$ is the dispersion parameter estimated from the global model, and can be calculated by dividing model's deviance by the number of residual degrees of freedom.

In calculation of QAIC, the number of model parameters is increased by 1 to account for estimating the overdispersion parameter. Without overdispersion, $\hat{c} = 1$ and QAIC is equal to AIC.

Note that glm does not compute maximum-likelihood estimates in models within the *quasi-* family. In case it is justified, and with a proper caution, a workaround could be used, by 'borrowing' the aic element from the corresponding 'non-quasi' family (see 'Example').

Author(s)

Kamil Bartoń

See Also

[AICc](#), [quasi](#) family used for models with over-dispersion

Examples

```
# Based on "example(predict.glm)", with one number changed to create
# overdispersion
budworm <- data.frame(
    ldose = rep(0:5, 2), sex = factor(rep(c("M", "F"), c(6, 6))),
    numdead = c(10, 4, 9, 12, 18, 20, 0, 2, 6, 10, 12, 16))
budworm$SF = cbind(numdead = budworm$numdead,
    numalive = 20 - budworm$numdead)

budworm.lg <- glm(SF ~ sex*ldose, data = budworm, family = binomial)
(chat <- deviance(budworm.lg) / df.residual(budworm.lg))

dredge(budworm.lg, rank = "QAIC", chat = chat)
dredge(budworm.lg, rank = "AIC")

## Not run:
# Ugly hacked constructor for quasibinomial family object, that allows for
# ML estimation
x.quasibinomial <- function(...) {
    res <- quasibinomial(...)
    res$aic <- binomial(...)$aic
    res
}
QAIC(update(budworm.lg, family = x.quasibinomial), chat = chat)

## End(Not run)
```

---

r.squaredLR   *Likelihood-ratio based pseudo-R-squared*

Description

Calculate a coefficient of determination based on the likelihood-ratio test (R^2_{LR}).

Usage

```
r.squaredLR(x, null = null.fit(x, TRUE))

null.fit(x, evaluate = FALSE, envir = environment(as.formula(formula(x))))
```

Arguments

x	a fitted model object.
null	a fitted *null* model, if not provided, a glm with only intercept and appropriate family will be used.
evaluate	If TRUE evaluate the fitted model object else return the call.
envir	the environment in which the *null* model is to be evaluated, defaults to the environment of the original model's formula.

Details

This statistic is is one of the several proposed pseudo-R-squared's for nonlinear regression models. It is based on an improvement from *null* (intercept only) model to the fitted model, and calculated as

$$R_{LR}^2 = 1 - \exp(-\frac{2}{n}(\log Lik(x) - \log Lik(0)))$$

where log*Lik(x)* and log*Lik(0)* are the log-likelihoods of the fitted and the *null* model respectively.

For OLS models the value is consistent with classical R^2. In some cases (e.g. in logistic regression), the maximum R_{LR}^2 is less than one. The modification proposed by Nagelkerke (1991) adjusts the R_{LR}^2 to achieve 1 at its maximum: $\bar{R}^2 = R_{LR}^2 / \max(R_{LR}^2)$ where $max(R_{LR}^2) = 1 - \exp(\frac{2}{n} \log Lik(0))$.

`null.fit` tries to guess the *null* model call (as a `glm`), given the provided fitted model object.

Value

`r.squaredLR` returns a value of R_{LR}^2, and the attribute `"adj.r.squared"` gives the Nagelkerke's modified statistic. Note that this is not the same as the classical 'adjusted R squared'.

`null.fit` returns the fitted *null* model object (if `evaluate = TRUE`) or an unevaluated call to fit a *null* model.

Note

R^2 is a useful goodness-of-fit measure as it has the interpretation of the proportion of the variance 'explained', but it performs poorly in model selection, and is not suitable for use in the same way as the information criterions.

References

Cox, D. R. and Snell, E. J. (1989) *The analysis of binary data*, 2nd ed. London, Chapman and Hall

Magee, L. (1990) R^2 measures based on Wald and likelihood ratio joint significance tests. *Amer. Stat.* 44: 250-253

Nagelkerke, N. J. D. (1991) A note on a general definition of the coefficient of determination. *Biometrika* 78: 691-692

See Also

summary.lm

---

subset.model.selection

Subsetting model selection table

Description

Return subsets of a model selection table returned by `dredge`.

Usage

```
## S3 method for class 'model.selection'
subset(x, subset, select, recalc.weights = TRUE, ...)
## S3 method for class 'model.selection'
x[i, j, recalc.weights = TRUE, ...]
```

Arguments

x	a model.selection object to be subsetted.
subset,select	logical expressions indicating columns and rows to keep. See subset.
i,j	indices specifying elements to extract.
recalc.weights	logical value specyfying whether Akaike weights should be normalized across the new set of models to sum to one.
...	further arguments passed to [.data.frame.

Value

A model.selection object containing only the selected models (rows). When columns are selected (arguments select or j are provided), a plain data.frame is returned.

Note

Unlike the method for data.frame, extracting with only one index (i.e. x[i]) will select rows rather than columns.

To select rows according to presence or absence of the variables (rather than their value), a pseudo-function has may be used, e.g. subset(x, has(a, !b)) will select rows with *a* **and** without *b* (this is equivalent to !is.na(a) & is.na(b)). has can take any number of arguments. Importantly, the has() notation cannot be used in the subset argument for dredge, where the variable names should be given directly, with the same effect.

Author(s)

Kamil Bartoń

See Also

dredge, subset and [.data.frame for subsetting and extracting from data.frames.

Examples

```
data(Cement)
fm1 <- lm(formula = y ~ X1 + X2 + X3 + X4, data = Cement)

# generate models where each variable is included only if the previous are
# included too, e.g. X2 only if X1 is there, and X3 only if X2 and X1 (this is
# not supposed to make sense in this case, just to demonstrate the syntax).
dredge(fm1, subset = (!X2 | X1) & (!X3 | X2) & (!X4 | X3))

# alternatively, generate "all possible" combinations
ms0 <- dredge(fm1)
# ...and afterwards select the subset of models
subset(ms0, (has(!X2) | has(X1)) & (has(!X3) | has(X2)) & (has(!X4) | has(X3)))
## Not run:
```

```
# this way the expression may be more clear
subset(ms0, has(X1, X2, X3, X4) |  has(X1, X2, X3) | has(X1, X2) | has(X1)
    | (df == 2))

## End(Not run)

# Different ways of finding a confidence set of models:
# delta(AIC) cutoff
subset(ms0, delta <= 4, recalc.weights = FALSE)
# cumulative sum of Akaike weights
subset(ms0, cumsum(weight) <= .95, recalc.weights = FALSE)
# relative likelihood
subset(ms0, (weight / weight[1]) > (1/8), recalc.weights = FALSE)
```

Weights   *Akaike weights*

Description

Calculate or extract normalized model likelihoods ('Akaike weights').

Usage

```
Weights(x)
```

Arguments

x   a numeric vector of information criterion values such as AIC, or objects returned
                by functions like AIC. There are also methods for extracting Akaike weights
                from a model.selection or averaging objects.

Value

a numeric vector of normalized likelihoods.

Author(s)

Kamil Bartoń

See Also

importance

weights, which extracts fitting weights from model objects.

Examples

```
data(Beetle)

fm1 <- glm(Prop ~ dose, data=Beetle, family=binomial)
fm2 <- update(fm1, . ~ . + I(dose^2))
fm3 <- update(fm1, . ~ log(dose))
fm4 <- update(fm3, . ~ . + I(log(dose)^2))
```

```
round(Weights(AICc(fm1, fm2, fm3, fm4)), 3)
```

Index